

Citibike Analysis: A Deeper Look Into New York's Bike Sharing Program

Author: Atharva Bhandarkar

Email: apb412@nyu.edu

Named after its chief sponsor, Citibank, Citibike is a privately owned public bike-sharing system in New York City. With 706 docking stations, 12000 bikes and attracting 163,000 annual subscribers in 2016, it is the largest bike sharing program in the United States.

Citibike has seen a meteoric rise in popularity over the last few years, growing from under a mere 6 million trips in 2013 to over 14 million in 2016. This massive surge in people's interest in and usage of the program made me curious about it and thus, the goal of this project is to provide a deeper insight about citibike usage in New York City.

The project will analyze data published by NYC Citibikes; first on a general level (some descriptive analysis), and gradually build on that by trying to answer certain, more thought-provoking questions.

The preliminary analysis will use data from a single month (or some small, fixed time period) and the following is a non-comprehensive list of questions it will answer:

- Average distance travelled by citibikes
- Gender and statistics of citibike users
- Most popular locations in NYC (start points, end points etc)

Once these questions are answered, the project will delve deeper and try to understand some more implicit 'trends'. What follows are some examples of these questions:

- How does usage change over the course of a day (i.e what are peak times)
- What does a weekly trend look like?

At this point we know who rides citibikes, how they ride them, when they ride them. But I wanted to know why. One would assume that they're faster than walking, but are they faster than driving? I used the Google Maps API to find out.

Data Report

Citibike publishes data about every single trip from 2012. As mentioned above, 2016 for example had 14 million trips; a datasize far too large (for computations sake). So I decided to focus on the September 2017 data.

The following is a link to where I obtained this zip file: <https://www.citibikenyc.com/system-data/> (<https://www.citibikenyc.com/system-data/>)

However because it is 330 mb, it made sense to download and use it from my laptop instead. In that case the path became:

```
'/Users/atharvabhandarkar/Downloads/citi_trips_sep2017.csv'
```

My Packages

```
In [1]: import pandas as pd
        from math import sin, cos, sqrt, atan2, radians
        import numpy as np
        import matplotlib.pyplot as plt
```

```
In [153]: path = '/Users/atharvabhandarkar/Downloads/citi_trips_sep2017.csv'
```

```
In [154]: df_full = pd.read_csv(path)
        df = df_full#.sample(n=100)
```

```
In [155]: df['bike_time'] = df['tripduration']/60 #creates new dataseries of trip t

        df = df.rename(columns={'start station latitude': 'start_lat',
                                'start station longitude': 'start_long',
                                'end station latitude': 'end_lat',
                                'end station longitude': 'end_long',
                                'start station name': 'start_station_name',
                                'end station name': 'end_station_name'})
```

Trip Distance Calculation

Using a trig function to calculate distances between start and end latitudes and longitudes of a trip

```
In [5]: def haversine(lon1, lat1, lon2, lat2):

    lon1, lat1, lon2, lat2 = map(np.radians, [df.start_long, df.start_lat,
                                                df.end_long, df.end_lat])

    dlon = lon2 - lon1
    dlat = lat2 - lat1

    a = np.sin(dlat/2.0)**2 + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2.0)**2
    c = 2 * np.arcsin(np.sqrt(a))
    km = 6367 * c
    return km
```

```
In [6]: for row in range(100): #doing range(100) because it is a 1.9million row dataframe
    #1.9mn times would be a unnecessarily time consuming
    df['dist_km'] = haversine(df.start_long.iloc[row],
                              df.start_lat.iloc[row],
                              df.end_long.iloc[row],
                              df.end_lat.iloc[row])
```

```
In [7]: df.head(2)
```

```
Out[7]:
```

	start station id	start_station_name	start_lat	start_long	end station id	end_station_name	end_lat	end_long
9- J1 19	3331	Riverside Dr & W 104 St	40.801343	-73.971146	3328	W 100 St & Manhattan Ave	40.795000	-73.9
9- J1 30	3101	N 12 St & Bedford Ave	40.720798	-73.954847	3100	Nassau Ave & Newell St	40.724813	-73.9

Some descriptive statistics about the distances travelled

```
In [10]: df.dist_km.mean()
```

```
Out[10]: 1.8863128567894845
```

```
In [12]: df.dist_km.describe() #1.87 million trips in Sept 2017, with average trip
```

```
Out[12]: count      1.878098e+06  
         mean      1.886313e+00  
         std       1.420290e+01  
         min       0.000000e+00  
         25%       8.755417e-01  
         50%       1.447624e+00  
         75%       2.423891e+00  
         max       8.662342e+03  
         Name: dist_km, dtype: float64
```

Age and Gender Distribution

```
In [60]: df['gender'] = df['gender'].astype(str) #converting the numerical field of  
                                                #so that they can be analyzed as c  
  
grouped_genders = df.groupby('gender').count() #grouping by gender  
sorted_gender = (grouped_genders.tripduration.sort_values(ascending = Fal  
sorted_gender
```

```
Out[60]: gender  
1      1218524  
2       444132  
0       215442  
         Name: tripduration, dtype: int64
```

There are 1,218,524 males, 444,132 females and 215,442 customers who aren't registered users hence their gender is unknown

```
In [37]: df['age'] = 2017-df['birth year']  
df.age.describe()
```

```
Out[37]: count      1.670630e+06  
         mean      3.746549e+01  
         std       1.171327e+01  
         min       1.600000e+01  
         25%       2.800000e+01  
         50%       3.500000e+01  
         75%       4.500000e+01  
         max       1.320000e+02  
         Name: age, dtype: float64
```

As the values in the series are floats (and attempts to convert them to integers failed due to NaN fields), summary statistics are displayed in scientific notation here.

- Average age of a citibike subscriber seems to be 37.46 years old.
- The mid 50% of the subscribers range between 28 and 45 years old.
- Furthermore, a standard deviation of 17 years is actually pretty high.

(There seems to be an error in birthyear input in the data field because according to the max age, there was a 132 year old rider in the month of Sept, which of course, seems highly unlikely.)


This next graph shows the cumulative frequency in male and female citi riders over different age ranges. It gives us a good idea as to how users are distributed in terms of gender and age. "

```

In [40]: fig, ax = plt.subplots()
df.groupby('gender').age.plot(ax = ax,
                              figsize = (15,10),
                              kind = 'hist',
                              alpha = 0.5,
                              title = 'Gender/Age Distribution of Citibike

ax.set_xlabel("Age")
ax.set_ylabel("Cumulative Frequency")
#ax.spines['left'].set_position(('outward', 10))
ax.spines['bottom'].set_position(('outward', 10))
ax.spines['top'].set_visible('False')
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
# ax.text(55, 150, 'Green: Women', fontsize = 12)
# ax.text(80, 140, 'Orange: Men', fontsize = 12)
legend = True
ax.set_xlim(16,80)
plt.show()

```



Most popular stations

```
In [41]: g = df.groupby('start_station_name').count() #sorting by start station name
g_sorted = g.tripduration.sort_values(ascending=False) #arranging in descending order
```

The 10 most popular start stations are:

```
In [42]: g_sorted.head(10)
```

```
Out[42]: start_station_name
Pershing Square North    17359
West St & Chambers St    13132
E 17 St & Broadway        12391
W 21 St & 6 Ave           11517
Broadway & E 22 St        11483
12 Ave & W 40 St          11356
8 Ave & W 31 St           11351
Broadway & E 14 St        10634
8 Ave & W 33 St           10226
W 38 St & 8 Ave           9901
Name: tripduration, dtype: int64
```

```
In [43]: g1 = df.groupby('end_station_name').count() #sorting by end station names
g_sorted1 = g1.tripduration.sort_values(ascending=False) #arranging in descending order
```

The 10 most popular end stations are:

```
In [44]: g_sorted1.head(10)
```

```
Out[44]: end_station_name
Pershing Square North    17056
West St & Chambers St    14917
E 17 St & Broadway        12704
Broadway & E 22 St        11830
W 21 St & 6 Ave           11668
8 Ave & W 31 St           11487
12 Ave & W 40 St          11361
Broadway & E 14 St        10617
W 20 St & 11 Ave          10263
8 Ave & W 33 St           9917
Name: tripduration, dtype: int64
```

What is interesting to observe is that almost all of these stations feature both in the most popular origination points and also destination points. And if the frequency for end station is higher, one can infer that there would have shortages of bikes at that dock at some point in the month.

For example Pershing Sq North has 17056 instances as destinations yet 17359 for instances of trip origination. This implies that the system

didn't replenish the bikes on its own and Citibike had to restock them themselves. Contrastingly the dock at West St and Chambers St seemingly maintained a healthy reserve of bikes with destinations outnumbering origination point data by 1785 bikes.

Given that the most popular origination point is Pershing Sq North, it begs the question, where are these people going? We will try to find out the most popular destinations from Pershing Sq North (PSN)

```
In [56]: PSN = df[df['start_station_name'] == 'Pershing Square North'] #create df w
PSN_grouped = PSN.groupby('end_station_name').count().reset_index() #grou
PSN_grouped.end_station_name.sort_values(ascending = False).head() #and s
```

```
Out[56]: 485          York St & Jay St
484      Wythe Ave & Metropolitan Ave
483          Wyckoff St & 3 Ave
482      Willoughby St & Fleet St
481      Willoughby Ave & Hall St
Name: end_station_name, dtype: object
```

Above are the top 5 destinations that people go to when they embark from Pershing Sq North

Most popular times; trends in usage

This section focuses on the timeseries aspect of the datafile. But the data itself is present in strings. So the following manipulation brings us to the desired state from where we can take off.

```
In [68]: starttime = df.starttime.tolist() #list of start time series values as st
stoptime = df.stoptime.tolist()

##remove seconds field because it created too many unique instances and w

for i in range(len(starttime)):
    starttime[i] = starttime[i][: -3]

for i in range(len(stoptime)):
    stoptime[i] = stoptime[i][: -3]

df['starttime'] = pd.to_datetime(starttime) #convert to datetime format
df['stoptime'] = pd.to_datetime(stoptime) #convert to datetime format
```

Usage of Citibikes on Sept 1

How usage fluctuated at various points in the day

- First we will group the trips by their start times (ie. 9.05am)
- Next we arrange them in descending order by the number of trips at each point in the day (ie. 43 trips at 9.05am)

Below is a plot of how citibike usage fluctuated throughout September 1st. The graph gives us a good idea of the peak times in the day.

```
In [71]: df0901 = df[(df['starttime'] > '2017-09-01 00:00') & (df['starttime'] <=
df0901.starttime.describe() # shows number of trips on the 1st of Sept
```

```
Out[71]: count          54877
unique          1390
top      2017-09-01 08:35:00
freq              109
first      2017-09-01 00:01:00
last      2017-09-01 23:59:00
Name: starttime, dtype: object
```


```
In [80]: df0901 = df[(df['starttime'] > '2017-09-01 00:00') & (df['starttime'] <=
sept_grouped = df0901.groupby('starttime').count() #sorts by starttimes
(sept_grouped.tripduration.sort_values(ascending = False)).head()
```

```
Out[80]: starttime
2017-09-01 08:35:00    109
2017-09-01 08:46:00    108
2017-09-01 08:52:00    105
2017-09-01 08:48:00    101
2017-09-01 09:16:00    100
Name: tripduration, dtype: int64
```

```
In [81]: fig, ax = plt.subplots(figsize = (18,10))
sept_grouped.tripduration.plot(kind = 'line')

ax.set_xlabel("Hour", fontsize = 20)
ax.set_ylabel("Number of Trips taken", fontsize = 20)
ax.spines['bottom'].set_position(('outward', 10))
ax.spines['top'].set_visible('False')
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
ax.set_title("September 1st Usage", fontsize = 20)

plt.show()
```



- We can see that the peak point in the day is 8.35am with 109 trips originating at that point.
- We can observe a general peak usage period around 8.30-9.15am which is logical due to the morning rush hour.
- There is further spikes in activity during the afternoon with 1.11pm and 2.12pm having 90 and 93 trips respectively.
- Furthermore, the evening rush hour period starts at around 6.18pm.

Exploring the week of Sept 1st to Sept 7th

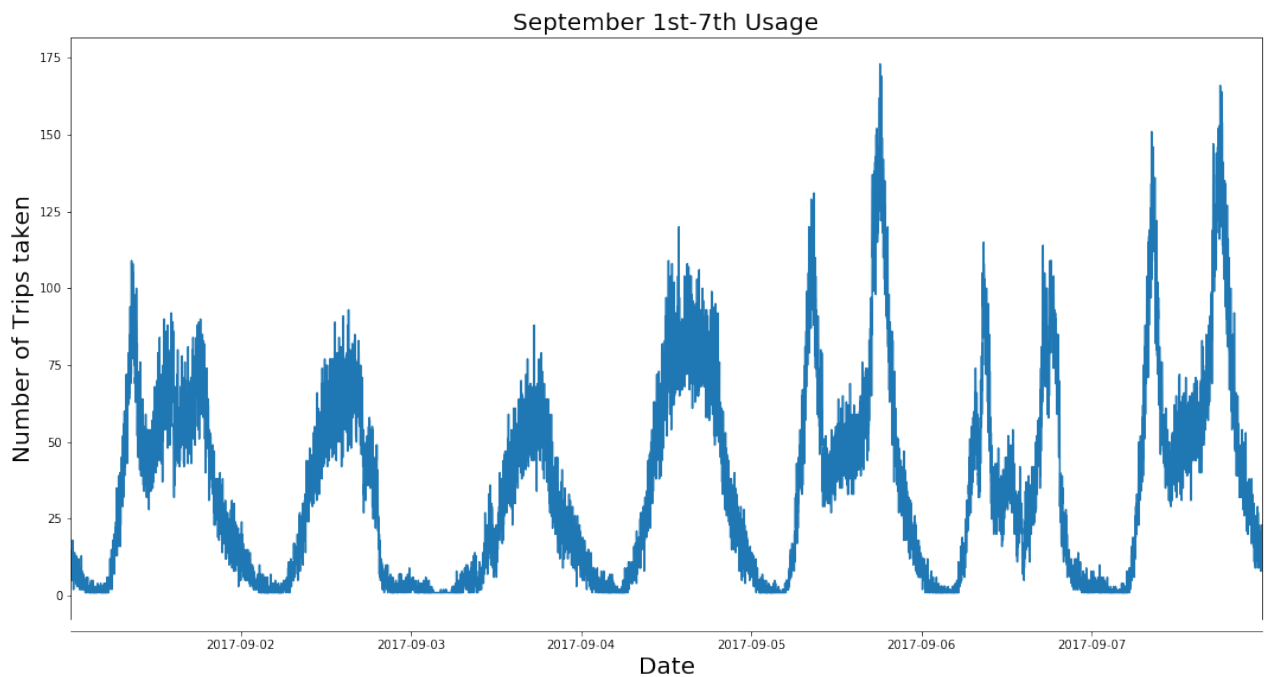
- We can then look at the trend throughout the week by changing the date parameters for the rows we import.
- Looking at the period between Sept 1st and Sept 7th, we can plot the following graph.
- We can clearly see a spike in the peak time activity on Sept 5th.

```
In [78]: df0901_07 = df[(df['starttime'] > '2017-09-01 00:00') & (df['starttime']

grouped = df0901_07.groupby('starttime').count() #sorts by starttimes
grouped2 = (grouped.tripduration.sort_values(ascending = False)).head()
fig, ax = plt.subplots(figsize = (18,10))
grouped.tripduration.plot(kind = 'line')

ax.set_xlabel("Date", fontsize = 20)
ax.set_ylabel("Number of Trips taken", fontsize = 20)
ax.spines['bottom'].set_position(('outward', 10))
ax.spines['top'].set_visible('False')
ax.yaxis.set_ticks_position('left')
ax.xaxis.set_ticks_position('bottom')
ax.set_title("September 1st-7th Usage", fontsize = 20)

plt.show()
```



```
In [79]: grouped2.head()
```

```
Out[79]: starttime
2017-09-05 18:08:00    173
2017-09-05 18:22:00    169
2017-09-05 18:13:00    167
2017-09-07 18:07:00    166
2017-09-07 18:18:00    164
Name: tripduration, dtype: int64
```

- A closer look at the peak times shows us that the peak was 6.08pm on Sept 5th with 173

trips.

- It is interesting to note that the peak period on Sept 5th was the evening rush hour unlike the morning rush hour we observed on Sept 1st.
- Also worth considering that Sept 4th is a Monday, and the days Sept 4th-7th (ie weekdays) show a gradual rise in activity over afternoon hours as compared to weekend days (Sept 2nd and 3rd), where they activity peaks in the afternoon.
- Interestingly, weekday peak usage periods are morning and evening rush hours (quite logically, of course) but what is strange is that peak usage on weekends is in the afternoon.
 - One could hypothesize that people are taking a leisurely bike ride around town or running errands in the afternoon.

Speed

Given distances and duration of our trips, we can easily calculate the average speeds of every trip.

```
In [124]: for row in range(100):  
           df0901['dist_km'] = haversine(df0901.start_long.iloc[row],  
                                         df0901.start_lat.iloc[row],  
                                         df0901.end_long.iloc[row],  
                                         df0901.end_long.iloc[row])  
  
           df0901['bike_time'] = df0901['tripduration']/60
```

```
In [125]: df0901['speed_kmh'] = df0901['dist_km']/(df0901['bike_time']/60)
```

```
In [126]: df0901.head()
```

```
Out[126]:
```

station_name	start_lat	start_long	end station id	end_station_name	end_lat	end_long	bikeid	u
Lafayette St & Jersey St	40.724305	-73.996010	3431	E 35 St & 3 Ave	40.746524	-73.977885	25413	Sul
it Ave & N 7 St	40.720368	-73.961651	3358	Garfield Pl & 8 Ave	40.671198	-73.974841	17584	Sul
W 106 St & msterdam Ave	40.800836	-73.966449	3343	W 107 St & Columbus Ave	40.799757	-73.962113	28581	Sul
' 20 St & 8 Ave	40.743453	-74.000040	388	W 26 St & 10 Ave	40.749718	-74.002950	30470	Sul
56 St & 10 Ave	40.768254	-73.988639	529	W 42 St & 8 Ave	40.757570	-73.990985	18150	Sul

Focus on the most popular citibike

This section will analyze the most popular bike in September using its bike-id

```
In [192]: df['bikeid'] = df['bikeid'].astype(str)
df['bikeid'].describe()
```

```
Out[192]: count      1878098
unique        11615
top           30458
freq           419
Name: bikeid, dtype: object
```

As we can see, although there have been 1,879,098 Citibike trips in the month of September, there are only 11,615 active bikes around the city. The most popular one seems to be bike number '30458' and has shown up in 419 different trips. This is the bike we will track.

```
In [194]: df30458 = df[df['bikeid']== '30458'] #create df with only rows that have
x = df30458.groupby('start_station_name').count()
x.sort_values(by= 'tripduration', ascending = False).head() # we use trip
#in descending
```

Out[194]:

	tripduration	starttime	stoptime	start station id	start_lat	start_long	end station id	end_s
start_station_name								
1 Ave & E 62 St	7	7	7	7	7	7	7	7
5 Ave & E 88 St	5	5	5	5	5	5	5	5
Carmine St & 6 Ave	5	5	5	5	5	5	5	5
Barrow St & Hudson St	4	4	4	4	4	4	4	4
W 18 St & 6 Ave	4	4	4	4	4	4	4	4

Above are the 5 most frequent points that bike 30458 originates from.

```
In [196]: df30458 = df[df['bikeid']== '30458'] #create df with only rows that have
x = df30458.groupby('end_station_name').count()
x.sort_values(by= 'tripduration', ascending = False).head() # we use trip
#in descending
```

Out[196]:

	tripduration	starttime	stoptime	start station id	start_station_name	start_lat	start_lo
end_station_name							
1 Ave & E 62 St	7	7	7	7	7	7	7
Carmine St & 6 Ave	5	5	5	5	5	5	5
5 Ave & E 88 St	5	5	5	5	5	5	5
W 18 St & 6 Ave	4	4	4	4	4	4	4
Greenwich Ave & 8 Ave	4	4	4	4	4	4	4

Above are the 5 most frequent locations that bike 30458 ends up at.

Interesting that 4 out of the 5 stations are the same as the top origination points. Could imply that many people just shuttle between those points and bike 30458 seems to be the most frequently used bike for this short trip.

Googlemaps Directions API

Using the API to figure out driving and walking times for the routes taken by the citibikes

Now for the biggest question; it it worth our while to take a city bike? To find out, we use to Google Maps API.

Sending the API the start and end station names for the trips in our dataset, the functions 'walking_time' and 'driving_time' will return the walking and driving times for the specific trips and enter these into our existing dataframe.

```
In [147]: # we use our start and end locations, the standard API format and an API
# we then query the API with this URL and extract a detailed set of direc
origin = df.start_station_name[0].replace(' ', '+') #remove whitespaces
destination = df.end_station_name[0].replace(' ', '+')
endpoint = 'https://maps.googleapis.com/maps/api/directions/json?'
api_key = 'AIzaSyCLNO5mol_LjqDuOkTKLBke4Q9de-6GVy4'

nav_request = 'origin={}&destination={}&key={}'.format(origin,destination)
request = endpoint + nav_request
response = urllib.request.urlopen(request).read()
directions = json.loads(response)

print (directions)
```

```
{'geocoded_waypoints': [{'geocoder_status': 'OK', 'place_id': 'ChIJ8aY
qYkH2wokRccUf0WRGIVE', 'types': ['route']}, {'geocoder_status': 'OK',
'place_id': 'ChIJfw68iyb2wokRjDwqmYkl9Mo', 'types': ['route']}], 'rout
es': [{'bounds': {'northeast': {'lat': 40.8162123, 'lng': -73.9620394}
, 'southwest': {'lat': 40.7960836, 'lng': -73.9700112}}, 'copyrights':
'Map data ©2017 Google', 'legs': [{'distance': {'text': '1.7 mi', 'val
ue': 2669}, 'duration': {'text': '8 mins', 'value': 494}, 'end_address
': 'W 100th St, New York, NY 10025, USA', 'end_location': {'lat': 40.7
960836, 'lng': -73.967017}, 'start_address': 'Riverside Dr, New York,
NY, USA', 'start_location': {'lat': 40.8162123, 'lng': -73.9620394}, '
steps': [{'distance': {'text': '1.0 mi', 'value': 1541}, 'duration': {
'text': '3 mins', 'value': 200}, 'end_location': {'lat': 40.8037, 'lng
': -73.9693882}, 'html_instructions': 'Head <b>southwest</b> on <b>Riv
erside Dr</b>', 'polyline': {'points': 'i|bxFvulbBMNRLNPLRL\\TD@\\N|EpB
vAn@lAh@f@TZJLDJBFB\\HVD\\BV@V?~AAX?H@L@N@RFFBVHLHxAbAJFFD|CpBpBrAtA~@
XPRXDF`@h@h@n@FFJHVR`@Xf@ZRLHFbB`APJNhrBz@~@b@VLXJz@^pAj@bAd@vAt@dBFaH
```

```
D'}, 'start_location': {'lat': 40.8162123, 'lng': -73.9620394}, 'travel_mode': 'DRIVING'}, {'distance': {'text': '0.1 mi', 'value': 165}, 'duration': {'text': '1 min', 'value': 45}, 'end_location': {'lat': 40.8029763, 'lng': -73.96767899999999}, 'html_instructions': 'Turn <b>left</b> onto <b>W 108th St</b>', 'maneuver': 'turn-left', 'polyline': {'points': 'cn`xftcnbMFSP]tBcH'}, 'start_location': {'lat': 40.8037, 'lng': -73.9693882}, 'travel_mode': 'DRIVING'}, {'distance': {'text': '0.4 mi', 'value': 677}, 'duration': {'text': '3 mins', 'value': 158}, 'end_location': {'lat': 40.7972537, 'lng': -73.9700112}, 'html_instructions': 'Turn <b>right</b> at the 1st cross street onto <b>Broadway</b>', 'maneuver': 'turn-right', 'polyline': {'points': 'si`xF~xmbM`A^TDPBL?J?LBPBBCLN@\\@vBNRAN@J@L@RDPDZHZFJDNDhBv@LDpBnALFJHdBfALHDDhBnAJF'}, 'start_location': {'lat': 40.8029763, 'lng': -73.96767899999999}, 'travel_mode': 'DRIVING'}, {'distance': {'text': '0.2 mi', 'value': 286}, 'duration': {'text': '2 mins', 'value': 91}, 'end_location': {'lat': 40.7960836, 'lng': -73.967017}, 'html_instructions': 'Turn <b>left</b> onto <b>W 100th St</b>', 'maneuver': 'turn-left', 'polyline': {'points': 'ye_xFpgnbMLa@bBoFJ]?Q?G?E?C@C`BmFFO'}, 'start_location': {'lat': 40.7972537, 'lng': -73.9700112}, 'travel_mode': 'DRIVING'}], 'traffic_speed_entry': [], 'via_waypoint': []}], 'overview_polyline': {'points': 'i|bxFvulbM\\b@d@Zb@VzF`ChFzB~@Vt@Hn@@xBaVBb@H^LzBzA~JvGdBzBlA~@hDrB`@TrD~A~DdBzCzAnBlAFSP]tBcH`A^f@HX?^FrCNtCPb@?~@NrA\\vB|@~BvAdC`BtBvApBqGJ]?Q?MjBeG'}, 'summary': 'Riverside Dr and Broadway', 'warnings': [], 'waypoint_order': []}], 'status': 'OK'}
```

```
In [148]: # the challenge here is to decypher the data dump google returns
# and to make sense of it so we can extract exactly what we need
# taking a closer glance at this chunk of text, we find that what we're r
# so we create a sublist called routes
routes = directions['routes']
routes
```

```
Out[148]: [{'bounds': {'northeast': {'lat': 40.8162123, 'lng': -73.9620394},
'southwest': {'lat': 40.7960836, 'lng': -73.9700112}},
'copyrights': 'Map data ©2017 Google',
'legs': [{'distance': {'text': '1.7 mi', 'value': 2669},
'duration': {'text': '8 mins', 'value': 494},
'end_address': 'W 100th St, New York, NY 10025, USA',
'end_location': {'lat': 40.7960836, 'lng': -73.967017},
'start_address': 'Riverside Dr, New York, NY, USA',
'start_location': {'lat': 40.8162123, 'lng': -73.9620394},
'steps': [{'distance': {'text': '1.0 mi', 'value': 1541},
'duration': {'text': '3 mins', 'value': 200},
'end_location': {'lat': 40.8037, 'lng': -73.9693882},
'html_instructions': 'Head <b>southwest</b> on <b>Riverside Dr</b>',
'polyline': {'points': 'i|bxFvulbMNRLNPLRL\\TD@\\N|EpBvAn@lAh@f@TZJLDJBFB\\HVD\\BV@V?~AAX?H@L@N@RFFBVLHxAbAJFFD|CpBpBrAtA~@XPRXDF`@h@h@n@FFJHVR`@Xf@ZRLHFbB`APJNHRbZ@~@b@VLXJz@^pAj@bAd@vAt@dBFaHD'},
'start_location': {'lat': 40.8162123, 'lng': -73.9620394},
```



```

        'travel_mode': 'DRIVING'},
        {'distance': {'text': '0.1 mi', 'value': 165},
         'duration': {'text': '1 min', 'value': 45},
         'end_location': {'lat': 40.8029763, 'lng': -73.96767899999999},
         'html_instructions': 'Turn <b>left</b> onto <b>W 108th St</b>',
         'maneuver': 'turn-left',
         'polyline': {'points': 'cn`xFtcnbMFSP]tBcH'},
         'start_location': {'lat': 40.8037, 'lng': -73.9693882},
         'travel_mode': 'DRIVING'},
        {'distance': {'text': '0.4 mi', 'value': 677},
         'duration': {'text': '3 mins', 'value': 158},
         'end_location': {'lat': 40.7972537, 'lng': -73.9700112},
         'html_instructions': 'Turn <b>right</b> at the 1st cross street
onto <b>Broadway</b>',
         'maneuver': 'turn-right',
         'polyline': {'points': 'si`xF~xmbM`A`TDPBL?J?LBPBbCLN@\\@vBNRAN@
J@L@RDPDZHZFJDNDhBv@LDpBnALFJHdBfALHDDhBnAJF'},
         'start_location': {'lat': 40.8029763, 'lng': -73.96767899999999}
    },
    {
        'travel_mode': 'DRIVING'},
        {'distance': {'text': '0.2 mi', 'value': 286},
         'duration': {'text': '2 mins', 'value': 91},
         'end_location': {'lat': 40.7960836, 'lng': -73.967017},
         'html_instructions': 'Turn <b>left</b> onto <b>W 100th St</b>',
         'maneuver': 'turn-left',
         'polyline': {'points': 'ye_xFpgnbMLa@bBoFJJ]?Q?G?E?C@C`BmFFO'},
         'start_location': {'lat': 40.7972537, 'lng': -73.9700112},
         'travel_mode': 'DRIVING'}],
    'traffic_speed_entry': [],
    'via_waypoint': [ ]],
    'overview_polyline': {'points': 'i|bxFvulbM\\b@d@Zb@VzF`ChFzB~@Vt@Hn
@@xBAVBb@H`LzBzA~JvGdBzBla~@hDrB`@TrD~A~DdBzCzAnBlAFSP]tBcH`A^f@HX?^Fr
CNtCPb@?~@NrA\\vB|@~BvAdC`BtBvApBqGJ]?Q?MjBeG'},
    'summary': 'Riverside Dr and Broadway',
    'warnings': [],
    'waypoint_order': [ ]]}

```

```

In [149]: # now its slightly less messy, but we still need to extract the exact time
          # turns out we're intersted in the values associated with 'legs'
          # so we create a list of 'legs' values

```

```

routes[0]['legs']

```

```

Out[149]: [{'distance': {'text': '1.7 mi', 'value': 2669},
            'duration': {'text': '8 mins', 'value': 494},
            'end_address': 'W 100th St, New York, NY 10025, USA',
            'end_location': {'lat': 40.7960836, 'lng': -73.967017},
            'start_address': 'Riverside Dr, New York, NY, USA',
            'start_location': {'lat': 40.8162123, 'lng': -73.9620394},
            'steps': [{'distance': {'text': '1.0 mi', 'value': 1541},

```

```

'duration': {'text': '3 mins', 'value': 200},
'end_location': {'lat': 40.8037, 'lng': -73.9693882},
'html_instructions': 'Head <b>southwest</b> on <b>Riverside Dr</b>'
',
'polyline': {'points': 'i|bxFvulbMNRLNPLRL\\TD@\\N|EpBvAn@lAh@f@TZ
JLDJBFB\\HVD\\BV@V?~AAX?H@L@N@RFFBVLHxAbAJFFD|CpBpBrAtA~@XPRXDF`@h@h@
n@FFJHVR`@Xf@ZRLHFbB`APJNHrBz@~@b@VLXJz@^pAj@bAd@vAt@dBFaHD'},
'start_location': {'lat': 40.8162123, 'lng': -73.9620394},
'travel_mode': 'DRIVING'},
{'distance': {'text': '0.1 mi', 'value': 165},
'duration': {'text': '1 min', 'value': 45},
'end_location': {'lat': 40.8029763, 'lng': -73.96767899999999},
'html_instructions': 'Turn <b>left</b> onto <b>W 108th St</b>',
'maneuver': 'turn-left',
'polyline': {'points': 'cn`xFTcnbMFSP]tBCH'},
'start_location': {'lat': 40.8037, 'lng': -73.9693882},
'travel_mode': 'DRIVING'},
{'distance': {'text': '0.4 mi', 'value': 677},
'duration': {'text': '3 mins', 'value': 158},
'end_location': {'lat': 40.7972537, 'lng': -73.9700112},
'html_instructions': 'Turn <b>right</b> at the 1st cross street on
to <b>Broadway</b>',
'maneuver': 'turn-right',
'polyline': {'points': 'si`xF~xmbM`A^TDPBL?J?LBPBbCLN@\\@vBNRAN@J@
L@RDPDZHZFJDNDhBv@LDpBnALFJHdBfALHDDhBnAJF'},
'start_location': {'lat': 40.8029763, 'lng': -73.96767899999999},
'travel_mode': 'DRIVING'},
{'distance': {'text': '0.2 mi', 'value': 286},
'duration': {'text': '2 mins', 'value': 91},
'end_location': {'lat': 40.7960836, 'lng': -73.967017},
'html_instructions': 'Turn <b>left</b> onto <b>W 100th St</b>',
'maneuver': 'turn-left',
'polyline': {'points': 'ye_xFpgnbMLa@bBoFJ]?Q?G?E?C@C`BmFFO'},
'start_location': {'lat': 40.7972537, 'lng': -73.9700112},
'travel_mode': 'DRIVING'}],
'traffic_speed_entry': [],
'via_waypoint': []}]

```

```

In [150]: #it is clear from this that the exact time is under duration
# we pin point the list by slicing it further and convert the seconds file
walk_secs = routes[0]['legs'][0]['duration']['value']/60
walk_secs

```

```
Out[150]: 8.233333333333333
```

```

In [151]: walk_mins=round(walk_secs,2) #rounding to 2 dp we get walk_mins
walk_mins

```

```
Out[151]: 8.23
```

```
In [133]: import urllib.request, json

# now we can create a function to automate this process a bit more
def walking_time(origin, destination): #give it start and end locations
    origin = df.start_station_name[row].replace(' ', '+') #remove whitesp
    destination = df.end_station_name[row].replace(' ', '+')
    endpoint = 'https://maps.googleapis.com/maps/api/directions/json?'
    api_key = 'AIzaSyCLN05mol_LjqDuOkTKLBke4Q9de-6GVy4'

    nav_request = 'origin={}&destination={}&key={}'.format(origin, destina
    request = endpoint + nav_request
    response = urllib.request.urlopen(request).read()
    directions = json.loads(response)
    routes = directions['routes']
    walk_secs = routes[0]['legs'][0]['duration']['value']/60
    walk_mins=round(walk_secs,2)

    return walk_mins
```

```
In [134]: walk_time= [] #creating a empty list of walk times where the values will
for row in range(len(df)):
    element = walking_time(df['start_station_name'][row], df['end_station
    walk_time.append(element) #add this value to the list
#walk_time
df['walk_time']= walk_time #create data series called walk_time
df.head()
```

Out[134]:

	tripduration	starttime	stoptime	start station id	start_station_name	start_lat	start_long	end station id
0	362	9/1/17 0:00	9/1/17 0:06	3331	Riverside Dr & W 104 St	40.801343	-73.971146	3328
1	188	9/1/17 0:00	9/1/17 0:03	3101	N 12 St & Bedford Ave	40.720798	-73.954847	3100
2	305	9/1/17 0:00	9/1/17 0:05	3140	1 Ave & E 78 St	40.771404	-73.953517	3141
3	223	9/1/17 0:00	9/1/17 0:04	236	St Marks Pl & 2 Ave	40.728419	-73.987140	473
4	758	9/1/17 0:01	9/1/17 0:13	3427	Lafayette St & Jersey St	40.724305	-73.996010	3431

```
In [152]: # we do the exact same thing for driving times
# the codes here are slightly different
#for example here instead of legs, we want legs[steps] and that retur
import urllib.request, json
## Driving Time Calculation
def driving_time(origin, destination):
    origin = df.start_station_name[row].replace(' ', '+')
    destination = df.end_station_name[row].replace(' ', '+')
    endpoint = 'https://maps.googleapis.com/maps/api/directions/json?'
    api_key = 'AIzaSyCLN05mol_LjqDuOkTKLBke4Q9de-6GVy4'

    nav_request = 'origin={}&destination={}&key={}'.format(origin, destination, api_key)
    request = endpoint + nav_request
    response = urllib.request.urlopen(request).read()
    directions = json.loads(response)
    routes = directions['routes']
    drive_secs = routes[0]['legs'][0]['steps'][0]['duration']['value']
    drive_mins = round(drive_secs/60,2)
    return drive_mins
```

```
In [136]: drive_time= [] #empty list of drive times
for row in range(len(df)):
    element = driving_time(df.start_station_name[row], df.end_station_name[row])
    drive_time.append(element) #append to empty list
df['drive_time'] = drive_time #create data series called walk_time
```

```
In [114]: df.head(10)
```

```
Out[114]:
```

start_station_name	start_lat	start_long	end_station_id	end_station_name	end_lat	end_long	bikeid
Riverside Dr & W 104 St	40.801343	-73.971146	3328	W 100 St & Manhattan Ave	40.795000	-73.964500	14530
N 12 St & Bedford Ave	40.720798	-73.954847	3100	Nassau Ave & Newell St	40.724813	-73.947526	15475
1 Ave & E 78 St	40.771404	-73.953517	3141	1 Ave & E 68 St	40.765005	-73.958185	30346
St Marks Pl & 2 Ave	40.728419	-73.987140	473	Rivington St & Chrystie St	40.721101	-73.991925	28056
Lafayette St & Jersey St	40.724305	-73.996010	3431	E 35 St & 3 Ave	40.746524	-73.977885	25413
Kent Ave & N 7 St	40.720368	-73.961651	3358	Garfield Pl & 8 Ave	40.671198	-73.974841	17584
W 106 St & Amsterdam Ave	40.800836	-73.966449	3343	W 107 St & Columbus Ave	40.799757	-73.962113	28581
W 20 St & 8 Ave	40.743453	-74.000040	388	W 26 St & 10 Ave	40.749718	-74.002950	30470
W 56 St & 10 Ave	40.768254	-73.988639	529	W 42 St & 8 Ave	40.757570	-73.990985	18150
Washington St & Gansevoort St	40.739323	-74.008119	358	Christopher St & Greenwich St	40.732916	-74.007114	28200

The above table neatly shows the Google Maps' values in the walk_time and drive_time columns.

It is worth mentioning that the reason I didn't do the entire dataframe is because as we know the dataframe is 1.87mn rows and the API has a per day limit of 100,000. But even then, the code itself takes time to run and 10 rows is enough to give an idea of the point we're trying to make.

So due to this stipulation and also simplicity sake we don't iterate through the entire dataframe.

Conclusion/ Final Words

As seen by the Google Maps API section, driving (unfortunately) does seem to win every race against the citibikes. But as predicted citibikes are obviously faster than walking. A further layer of complexity to this investigation could have been analyzing how this changes during peak times. For example would driving still be the fast option during peak times of the day?

In []: